

mimix: Reference-based imputation of missing data

by Ian R White, Kevin McGrath, Matteo Quartagno, Suzie M Cro and James Carpenter

Abstract In some randomised trials, outcome data are missing after participant discontinuation of allocated treatment. Reference-based imputation is a multiple imputation technique for quantitative outcome data in this setting: it constructs the imputation model using parameters from a reference group. We present and describe an R package, *mimix*, for performing reference-based imputation. The package also includes a causal model variant and imputation of data missing before discontinuation of allocated treatment. We compare its implementation with that in SAS and Stata.

Introduction

Missing data are a challenge for many analyses. This article tackles the specific issue of a randomised trial with a repeatedly measured quantitative outcome, where participants who discontinue their randomised treatment are not followed up thereafter and hence have missing outcome data. We assume that the aim is to estimate the effect of treatment on the actual outcomes of the participants, whether observed or not: this has been termed a “de facto” estimand (Carpenter et al., 2013) or a “treatment policy” estimand (International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use, 2019). In this setting, an analysis under the commonly used missing at random (MAR) assumption would assume that the missing (post-treatment) outcomes are comparable (conditional on observed data) with the observed (on-treatment) outcomes, and would therefore estimate the effect of treatment on the outcomes of the participants if treatment was never discontinued: this has been termed a “de jure” estimand (Carpenter et al., 2013) or a “hypothetical” estimand (International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use, 2019).

Carpenter et al. (2013) proposed that missing data after discontinuation of allocated treatment could be imputed by assuming that post-discontinuation outcomes behave, in some sense, like the outcomes in a reference group. For example, if participants who have discontinued their allocated treatment are likely to be receiving similar treatment to participants who were allocated to control treatment, then the control group would be the reference group. Carpenter et al. (2013) proposed imputing the missing data from a joint multivariate Normal (MVN) distribution for the complete (observed and unobserved) data, and proposed five ways to construct this joint distribution: “jump to reference” (J2R), “copy reference” (CR), “copy increments in reference” (CIR), “last mean carried forward” (LMCF) and “missing at random” (MAR). All approaches start by fitting a MVN distribution to the data from each arm. As an example, the J2R joint distribution for a participant in a specific arm takes the means for that arm up to the point of treatment discontinuation and the means for the reference arm afterwards.

The implicit assumptions behind this approach were explored by White et al. (2020), who proposed a causal model in which the treatment effect after discontinuation is a specified fraction of the treatment effect at the point of discontinuation. They showed that J2R, CR and CIR are the special cases of this causal model in which the treatment effect disappears, decays or is maintained after treatment discontinuation, while LMCF and MAR are not special cases of the causal model.

The RBI methods of Carpenter et al. (2013) were implemented in SAS in what have become known as “the five macros” and are available on the web page (on www.missingdata.org.uk) of the DIA working group for missing data. These macros also provide for “delta adjustment” in which imputed values are modified by a user-specified amount (Ratitch et al., 2013): this is useful in performing sensitivity analysis, since the RBI methods make a number of untestable assumptions (Leacy et al., 2017). The methods were then implemented in Stata by Cro et al. (2016). A practical guide to their use is given by Cro et al. (2020).

We have implemented the RBI methods in a new R package, which includes the a full implementation of the causal model. The causal model was previously specified for a two-arm trial and has been extended here for a multi-arm trial. The aim of this paper is to describe our implementation of *mimix* in R, to illustrate its use, and to describe its novel features by comparison with the implementations in SAS and Stata.

Reference-based imputation

Setting

Following [White et al. \(2020\)](#), we assume that quantitative outcome measurements are scheduled at baseline and at T occasions after randomisation. Let Z be the random variable for the participant's randomised treatment arm, where $Z = z$ denotes an arm in which we want to impute missing values, and $Z = r$ denotes the reference arm. Let Y_t be the random variable for the participant's outcome at visit $t = 0, \dots, T$. It is convenient to define row vectors $\mathbf{Y}_{\leq t} = (Y_0, \dots, Y_t)$ containing all outcomes up to and including visit t ; $\mathbf{Y}_{> t} = (Y_{t+1}, \dots, Y_T)$ containing all outcomes after visit t ; and $\mathbf{Y} = (Y_0, \dots, Y_T)$, containing all outcomes. Let D be the random variable for the participant's last visit prior to discontinuing treatment, so $D = 0, \dots, T$. Y_t is observable for all t but only observed for $t \leq D$, because we assume no off-treatment data. We aim to impute the unobserved values of Y_t for $t > D$: we stress that these are the outcomes that existed but were unobserved, not the outcomes that would have existed if treatment had been continued.

Reference-based imputation

[Carpenter et al. \(2013\)](#) proposed a generic MI algorithm for this setting:

1. For each treatment arm z , fit a multivariate normal model to all observed data, using a Bayesian approach with an improper prior and assuming MAR. The model has unstructured mean μ_z and unstructured variance-covariance matrix Σ_z .
2. For each treatment arm z , draw a mean vector μ_z^* and variance-covariance matrix Σ_z^* from the posterior distribution.
3. For each treatment arm z and each observed treatment discontinuation visit t , use the drawn values to build a hypothetical joint distribution of the outcomes $\mathbf{Y}_{\leq t}$ up to visit t and the outcomes $\mathbf{Y}_{> t}$ after visit t , using one of the methods described below. Thus a MVN distribution is built for $\mathbf{Y}|Z = z, D = t$. We define the means $\mu_{z, \leq t}^*$, $\mu_{z, > t}^*$ to be the parts of μ_z^* up to and after visit t respectively.

Five methods are mainly distinguished by their choice of mean:

- (a) Jump to reference (J2R): mean = $(\mu_{z, \leq t}^*, \mu_{r, > t}^*)$.
- (b) Copy reference (CR): mean = μ_r^* .
- (c) Copy increments in reference (CIR): mean = $(\mu_{z, \leq t}^*, \{\mu_{z, t}^* - \mu_{r, t}^*\} \mathbf{e}_{T-t} + \mu_{r, > t}^*)$ where \mathbf{e}_p is a row vector $(1, \dots, 1)$ of length p .
- (d) Missing at random (MAR): mean = μ_z^* .
- (e) Last mean carried forward (LMCF): mean = $(\mu_{z, \leq t}^*, \mu_{z, t}^* \mathbf{e}_{T-t})$.

The variance matrices are constructed so that the regression coefficient matrix and conditional variance matrix of the potential outcomes after visit t , given those before visit t , are taken from arm z for MAR and LMCF, and from arm r for J2R, CIR and CR [Carpenter et al. \(2013\)](#); [White et al. \(2020\)](#).

4. For each treatment arm z and each observed treatment discontinuation visit t , construct the imputation distribution of $\mathbf{Y}_{> t}$ given $\mathbf{Y}_{\leq t}$. Sample $\mathbf{Y}_{> t}$ from this conditional distribution, to create a "completed" data set.
5. Repeat steps 2-4 m times, resulting in m imputed data sets.
6. Analyse each imputed data set and combine the results using Rubin's rules ([Rubin, 1987](#)).

Causal model

The causal model has previously been stated for two arms only. Here we extend its statement to the multi-arm case. We define the potential outcome $Y_t(z, s)$ at visit t as the outcome that would have been observable if, possibly contrary to fact, a participant received active treatment z for s periods only. In particular, $Y_t(z, 0)$ is the potential outcome if never treated: it is the same for all z and is written $Y_t(0)$. Similarly $Y_t(z, T)$ is the potential outcome if always treated with z . We let $\mu_t(z, s) = \mathbb{E}[Y_t(z, s)]$, the mean of the potential outcome at visit t if active treatment z were received for s periods only. We define row vectors $\mathbf{Y}_{\leq t}(z, s)$, $\mathbf{Y}_{> t}(z, s)$, $\mathbf{Y}(z, s)$, $\mu_{\leq t}(z, s)$, $\mu_{> t}(z, s)$ and $\mu(z, s)$ as before.

[Omit?:] The variance-covariance matrix of the potential outcomes is $\Sigma(s) = \text{var}(Y(s))$ with sub-matrices $\Sigma_{\leq t \leq t}(s)$, $\Sigma_{> t > t}(s)$ and $\Sigma_{> t \leq t}(s)$. We define the matrix of regression coefficients of potential outcomes after visit t on those up to visit t as $\beta_t(s) = \Sigma_{> t \leq t}(s) \Sigma_{\leq t \leq t}(s)^{-1}$, and the residual variance of the potential outcomes after visit t given those up to visit t as $\Omega_t(s) = \Sigma_{> t \leq t}(s) \Sigma_{\leq t \leq t}(s)^{-1} \Sigma_{> t \leq t}(s)^T$.

The key model assumption describes how the maintained effect of treatment after discontinuation relates to the effect of treatment before discontinuation:

$$\mu_{>t}(z, t) - \mu_{>t}(0) = K_{z,t} \{ \mu_{\leq t}(z, t) - \mu_{\leq t}(0) \} \quad (1)$$

where $K_{z,t}$ is a $(T - t) \times (t + 1)$ matrix of sensitivity parameters: it is not identified by the data and must be specified by the user. In practice we make a choice of $K_{z,t}$ determined by just two parameters (k_0, k_1) giving the simplified model for each $u > t$:

$$\mu_u(z, t) - \mu_u(0) = k_0 k_1^{v_u - v_t} \{ \mu_t(z, t) - \mu_t(0) \} \quad (2)$$

where v_t is the time (on a suitable scale) of visit t .

[At present (k_0, k_1) are fixed but it may be possible to let them depend on arm. This makes sense for a trial with more than 2 arms. Would probably code them as a variable in data set so could let them vary by any individual factors e.g. reason for missing data]

Estimation involves other assumptions which make explicit the ideas of [Carpenter et al. \(2013\)](#): that randomisation is independent of all potential outcomes; that step 1 of the RBI algorithm in the active arms estimates the distribution of $Y(z, s)$, and in the reference arm estimates the distribution of $Y(0)$; that the conditional distributions follow linear regressions (which is true if the joint distribution is MVN); and that treatment discontinuation is unaffected by future potential partly-treated outcomes.

These give the mean of the imputation model:

$$\begin{aligned} E[Y_{>t}(z, t) | Z = z, Y_{\leq t}, D = t] \\ = \beta_t(z, t) \{ Y_{\leq t} - \mu_{\leq t}(z, t) \} + K_{z,t} \{ \mu_{\leq t}(z, t) - \mu_{\leq t}(0) \} + \mu_{>t}(0). \end{aligned} \quad (3)$$

[Again we can handle variances... how much to say?]

Interim missing values

Interim missing values are values that are missing while the individual remains on treatment and are identified by being followed by later observed values. In the RBI algorithm and causal model, applying the constructed imputation distribution to missing values in $Y_{\leq t}$ implies imputation under MAR. We consider MAR imputation of interim missing values to be the procedure that is most likely to be appropriate in practice, since interim missing values that occur while on treatment are unlikely to follow the same pattern as post-discontinuation missing values that occur when off treatment. This differs from the behaviour of the Stata package ([Cro et al., 2016](#)) which by default imputes interim missing values using the same procedure as post-discontinuation values, with MAR as an option.

Individuals with no observed values

[TEXT FOR DISCUSSION – NEED TO WORK OUT HOW TO HANDLE THESE DISCREPANCIES]

Individuals with no observed values are easily imputed under MAR, but for the other methods they are currently a problem that we need to resolve.

We first consider J2R and CR. These methods should construct the imputation distribution wholly from the reference arm. They therefore correspond to MAR imputation if the wholly unobserved individual was in the reference arm. Both the Stata and R programs do this correctly for an individual in the reference arm, but for an individual in a non-reference arms, J2R and CR do not agree with each other.

We next consider CIR and LMCF. These are defined in terms of the last observed value, so they are formally undefined. The Stata program imputes CIR the same as J2R, while LMCF differs from all other methods, except that it agrees with MAR at the first visit. The R program imputes CIR and J2R differently when in the non-reference arm.

Delta adjustment

Delta adjustment allows imputations to differ systematically from RBI methods and so provides a framework for performing sensitivity analysis around RBI assumptions. We follow the approach of the five macros. Delta adjustment provides an increment which is added on to all values imputed after treatment discontinuation, but not to interim missing values. It is specified by two user-specified T -dimensional row vectors \mathbf{a} and \mathbf{b} , where a_u is a shift for time u , and b_w is a scaling multiplier that controls how the user-defined shift for time u is applied at time $u + w$. Formally, any imputed value $Y_{z,t}^*$ is replaced by $Y_{z,t}^* + \sum_{u=t+1}^{u=T} a_u b_{u-t}$. This means that elements of \mathbf{a} are cumulated after treatment discontinuation.

For example, consider an individual who discontinued treatment at the 2nd time point. Under the setting $b_w = 1$ for all w , we take the vector \mathbf{a} starting at the 3rd element and add the cumulative sums to the imputed values. Modifying \mathbf{b} modifies this behaviour, so that the vector of \mathbf{a} starting at the 3rd time point is multiplied elementwise by the vector \mathbf{b} . A common increment of 3 at all time points after treatment discontinuation would be achieved by setting $\mathbf{a} = (3, 3, 3, \dots)$ and $\mathbf{b} = (1, 0, 0, \dots)$. An increment of 3 at the time point immediately after discontinuation that is halved at each subsequent time point is specified by $\mathbf{a} = (3, 3, 3, \dots)$ and $\mathbf{b} = (1, -1/2, -1/4, \dots)$.

Delta adjustment can apply either with RBI or the causal method and is always applied after imputation.

Covariates

Covariates are used to improve the plausibility of MAR assumptions and to improve the precision of imputed values (White et al., 2011). In `mimix`, they are included in the multivariate normal model fitting as part of the variable Y_0 .

Missing values in covariates are not allowed in `mimix`. For these missing values in the randomised setting, single imputation methods are preferred, so mean imputation or the missing indicator method should be used before calling `mimix` (Sullivan et al., 2018).

Package

Arguments

Option	Description
<i>Data options</i>	
data	Dataset in long format, i.e. with one record per individual per time point.
covar	Covariates measured at or before baseline. Must be complete (no missing values).
depvar	Dependent (outcome) variable.
treatvar	Treatment group, coded 1,2,...
idvar	Participant identifier.
timevar	Time point for repeated measures.
<i>Method options</i>	
method	Reference-based imputation method: must be one of J2R, CIR, CR, LMCF, MAR, Causal.
methodvar	Alternative to method option: variable in dataset specifying the imputation method for each individual.
reference	Reference group for J2R, CIR, CR and Causal methods.
referencevar	Alternative to reference option: variable in dataset specifying the reference group for each individual.
<i>Causal method options</i>	
K0	Causal constant for use with Causal method. The treatment effect assumed at post-discontinuation times is K0 times the treatment effect at the time of discontinuation; if $K1 \neq 1$, this is multiplied by a decaying term.
K1	Causal constant for use with Causal method. The treatment effect assumed at post-discontinuation times, implied by K0, decays exponentially by a factor K1 per time unit.
<i>Delta method options</i>	
delta	Optional vector <i>a</i> of values to add onto imputed values as described above under “Delta adjustment”. Length must equal number of time points.
dlag	Optional vector <i>b</i> of values to modify the values in delta. Length must equal number of time points.
<i>Computation options</i>	
M	Number of imputations to be created.
seed	Seed value. Specify this so that a new run of the command will give the same imputed values.
<i>MCMC options</i>	
prior	Prior for the variance-covariance matrix when fitting multivariate normal distributions: Jeffreys (default), uniform or ridge.
burnin	Number of burn-in iterations when fitting multivariate normal distributions.
bbetween	Number of iterations between imputed data sets when fitting multivariate normal distributions.
mle	Logical option to use maximum likelihood parameter estimates instead of MCMC draw parameters. [WE'RE EXPLORING WHETHER WE CAN HIDE THIS OPTION, WHICH WAS ONLY WRITTEN TO PLEASE JONATHAN BARTLETT]

Algorithm

The program works through the following steps.

1. Set up a summary table based on treatment arm and missing data pattern (i.e. which timepoints are unobserved).
2. Fit a multivariate normal distribution to each treatment arm using MCMC methods in package norm2 (Schafer, 2021).
3. Impute all interim missing values under a MAR assumption, looping over treatments and patterns.
4. Impute post-discontinuation missing values under the user-specified assumption, looping over treatments and patterns. If methodvar and/or refvar is specified then the loop is also over the values of these.
5. Perform delta adjustment if specified.
6. Repeat steps 2-5 M times and form into a single data frame.

The baseline value of the outcome could be handled as an outcome or as a covariate. [Which do we recommend?]

The program is based on Suzie Cro's Stata program `mimix`.

Outputs

The M imputed data sets are output concatenated as one large data frame, comprising the original unimputed dataset concatenated with all the imputed data sets. [CURRENTLY IN WIDE FORMAT, UNLIKE THE INPUT DATA AND UNLIKE STATA'S OUTPUT]

The user can use the `as.mids()` function in the `mice` package to convert the output data to `mids` data type and hence to perform analysis using Rubin's rules.

Example: asthma trial data

These data are taken from [Cro et al. \(2016\)](#) and represent 2 arms of a 5-arm trial. They contain 183 patients with asthma who were randomised to active treatment or placebo. The primary outcome is FEV at 12 weeks, and FEV is also measured at baseline and at 2, 4, and 8 weeks of follow-up. The data contain 732 records (one per patient per follow-up visit) and 5 variables: `id`, a patient identifier; `time`, the visit number; `treat`, the randomised group (1=placebo, 2=active); `base`, the baseline FEV, to be regarded as a covariate; and `fev`, the outcome variable.

The data are missing after patients discontinue their randomised treatment. A treatment policy estimand is therefore estimated using reference-based imputation, with the placebo arm as reference. For example, J2R is implemented as follows

```
asthmaJ2R <- mimix(data=asthma, covar='base', depvar=fev, treatvar=treat, idvar=id,
  timevar=time, method='J2R', reference=1, M=5, seed=101, prior=ridge, burnin=1000)
```

The primary outcome can be analysed using an analysis of covariance by

```
library(mice)
fit <- with(data= as.mids(asthmaJ2R), lm(fev.12 ~ treat + base))
summary(pool(fit))
```

Results are as follows:

	term	estimate	std.error	statistic	df	p.value
1	(Intercept)	0.5821281	0.21097369	2.759245	19.15041	1.242142e-02
2	treat.L	0.1585223	0.07119548	2.226578	81.64041	2.873017e-02
3	base	0.6843152	0.09180436	7.454060	28.90698	3.315976e-08

Comparisons with other packages

[TO DISCUSS – TEXT, TABLE OR BOTH?]

Comparison with Stata

This `mimix` is based on the Stata version and has similar functionality while adding the causal method and delta adjustment. As with the Stata version the input data requires the longitudinal input data in long format with one record per individual at each timepoint.

The program differs in how interim missing cases – those cases which have a missing measurement at a timepoint previous to a later observed measurement – are treated. Under Stata by default, the interim missing are treated the same as for the post-discontinuation missing unless the `interim` option is explicitly used. Here the interims are treated as under MAR, the post-discontinuations then imputed under the specified method. There is no `interim` option as there is in Stata.

Unlike Stata an option is supplied whereby the prior used in the MCMC draws can be changed from the default `jeffreys` (as in Stata) to either the `ridge` or `uniform`.

Comparison with SAS

Whilst this program is based on the Stata program, the latter is an adaptation of the SAS macro `miwithd`, written by James Roger, subsequently updated to the `Five_Macros` suite of macros. This program uses the same approach for the delta adjustment as described in the `Five_Macros`.

In comparing outputs from our program with the `Five_Macros` it is to be noted that interaction between treatment and covariates is not allowed in the SAS macros, and comparisons are only valid for example in testing the Causal model by specifically not using the `covbytime` and `catcovbytime` options in the `Five_Macros`.

Not using these options also means that the LMCF method can be compared with either ALMCF or OLMCF in the `Five_Macros`. [I STILL HAVE TO WORK OUT WHICH WE HAVE IMPLEMENTED.]

When there is no observed data (common in the acupuncture data) the first mean is used in Stata, a warning is given in the `Five_Macros`.

summary table

	Five macros in SAS	mimix in Stata	mimix in R
<i>Functionality</i>			
RBI	yes	yes	yes
Causal model	yes	no	yes
Delta method	yes	no	yes
<i>Options</i>			
Interim missing	MAR	RBI or MAR	MAR
MCMC options	seed and thinning (=burnin)	seed, burnin and thinning	seed, burnin, thin- ning and prior
Covariates interact- ing with treatment	never	always	always
Covariates interact- ing with time	optional	always	always
LMCF	several flavours: ALMCF, OLMCF, AFCMCF, OFCMCF	ALMCF only?	ALMCF only?
No observed data	Warning	First mean is used	?

Limitations and discussion: sketch only so far

- Flexible package
- doesn't cover other outcome types; e.g. methods have been proposed for recurrent events (Keene et al., 2014)
- The suitability of the Rubin's rules standard errors has been debated (Seaman et al., 2014; Carpenter et al., 2014; Cro et al., 2019)
- Trials should be designed consistently with their estimand, so if the treatment-policy estimand is of interest then outcomes should be collected after treatment discontinuation. Analysis options for this setting are still in development: options include including treatment discontinuation time in the model and imputing under MAR [ref Tom at GSK]; using RBI or causal model and reserving the post-discontinuation data for model checking; or using RBI or causal model and using the post-discontinuation data to estimate model parameters such as $K_{z,t}$.
- We didn't code the *RBI alternative* approach of White et al. (2020) that instead uses $\beta_t(T)$ and $\Omega_t(T)$ for all RBI methods.

Bibliography

J. R. Carpenter, J. H. Roger, and M. G. Kenward. Analysis of longitudinal trials with protocol deviation: a framework for relevant, accessible assumptions, and inference via multiple imputation. *Journal of Biopharmaceutical Statistics*, 23(3):1352–71, 2013. ISSN 1520-5711. doi: 10.1080/10543406.2013.834911. URL <http://www.tandfonline.com/doi/abs/10.1080/10543406.2013.834911>. [p1, 2, 3]

J. R. Carpenter, J. H. Roger, S. Cro, and M. G. Kenward. Response to comments by Seaman et al. on “Analysis of longitudinal trials with protocol deviation: a framework

- for relevant, accessible assumptions, and inference via multiple imputation", *Journal of Biopharmaceutical Statistics* 23:1352-1371. *Journal of Biopharmaceutical Statistics*, 24 (6):1363-1369, nov 2014. ISSN 1054-3406. doi: 10.1080/10543406.2014.960085. URL <http://www.tandfonline.com/doi/abs/10.1080/10543406.2014.960085?journalCode=lbps20{#}.VcxWpmd7R8Ehttp://www.tandfonline.com/doi/abs/10.1080/10543406.2014.960085>. [p7]
- S. Cro, T. P. Morris, M. G. Kenward, and J. R. Carpenter. Reference-based sensitivity analysis via multiple imputation for longitudinal trials with protocol deviation. *Stata Journal*, 16(2):443-463, 2016. [p1, 3, 6]
- S. Cro, J. R. Carpenter, and M. G. Kenward. Information-anchored sensitivity analysis: theory and application. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*, 182(2):623-645, feb 2019. ISSN 09641998. doi: 10.1111/rssa.12423. URL <http://doi.wiley.com/10.1111/rssa.12423>. [p7]
- S. Cro, T. P. Morris, M. G. Kenward, and J. R. Carpenter. Sensitivity analysis for clinical trials with missing continuous outcome data using controlled multiple imputation: A practical guide. *Statistics in Medicine*, 39(21):2815-2842, sep 2020. ISSN 0277-6715. doi: 10.1002/sim.8569. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.8569>. [p1]
- International Council for Harmonisation of Technical Requirements for Pharmaceuticals for Human Use. Addendum on estimands and sensitivity analysis in clinical trials to the guideline on statistical principles for clinical trials E9(R1). Technical report, 2019. URL <https://www.ich.org/page/efficacy-guidelines{#}9-2>. [p1]
- O. N. Keene, J. H. Roger, B. F. Hartley, and M. G. Kenward. Missing data sensitivity analysis for recurrent event data using controlled imputation. *Pharmaceutical Statistics*, 13(4):258-264, jul 2014. ISSN 15391604. doi: 10.1002/pst.1624. URL <http://doi.wiley.com/10.1002/pst.1624>. [p7]
- F. P. Leacy, S. Floyd, T. A. Yates, and I. R. White. Analyses of Sensitivity to the Missing-at-Random Assumption Using Multiple Imputation With Delta Adjustment: Application to a Tuberculosis/HIV Prevalence Survey With Incomplete HIV-Status Data. *American Journal of Epidemiology*, 185(4):304-315, jan 2017. ISSN 0002-9262. doi: 10.1093/aje/kww107. URL <https://academic.oup.com/aje/article-lookup/doi/10.1093/aje/kww107http://aje.oxfordjournals.org/lookup/doi/10.1093/aje/kww107>. [p1]
- B. Ratitch, M. O'Kelly, and R. Tosiello. Missing data in clinical trials: from clinical assumptions to statistical analysis using pattern mixture models. *Pharmaceutical Statistics*, 12:337-347, 2013. URL <http://onlinelibrary.wiley.com/doi/10.1002/pst.1549/abstract>. [p1]
- D. B. Rubin. *Multiple Imputation for Nonresponse in Surveys*. John Wiley and Sons, New York, 1987. [p2]
- J. L. Schafer. CRAN - Package norm2, 2021. URL <https://cran.r-project.org/web/packages/norm2/index.html>. [p5]
- S. Seaman, F. Leacy, and I. R. White. Re: Analysis of Longitudinal Trials with Protocol Deviations — a Framework for Relevant, Accessible Assumptions, and Inference via Multiple Imputation (Carpenter, Roger and Kenward, *Journal of Biopharmaceutical Statistics* 2013;23:1352-1371). *Journal of Biopharmaceutical Statistics*, 24:1358-1362, 2014. URL <http://www.tandfonline.com/doi/full/10.1080/10543406.2014.928306>. [p7]
- T. R. Sullivan, I. R. White, A. B. Salter, P. Ryan, and K. J. Lee. Should multiple imputation be the method of choice for handling missing data in randomized trials? *Statistical Methods in Medical Research*, 27(9):2610-2626, sep 2018. ISSN 0962-2802. doi: 10.1177/0962280216683570. URL <http://journals.sagepub.com/doi/10.1177/0962280216683570>. [p4]
- I. R. White, A. Wood, and P. Royston. Tutorial in biostatistics: Multiple imputation using chained equations: issues and guidance for practice. *Stat Med*, 30:377-399, 2011. URL <http://onlinelibrary.wiley.com/doi/10.1002/sim.4067/abstract>. [p4]
- I. R. White, R. Joseph, and N. Best. A causal modelling framework for reference-based imputation and tipping point analysis in clinical trials with quantitative outcome. *Journal of Biopharmaceutical Statistics*, 30(2):334-350, mar 2020. ISSN 1054-3406. doi: 10.1080/10543406.2019.1684308. URL <https://www.tandfonline.com/doi/full/10.1080/10543406.2019.1684308>. [p1, 2, 7]

Ian R White
MRC Clinical Trials Unit at UCL
90 High Holborn, 2nd Floor, London WC1V 6LJ
UK
ORCID: 0000-0002-6718-7661
ian.white@ucl.ac.uk

Kevin McGrath
MRC Clinical Trials Unit at UCL
90 High Holborn, 2nd Floor, London WC1V 6LJ
UK
(ORCID if desired)
kevin.mcgrath@ucl.ac.uk

Matteo Quartagno
MRC Clinical Trials Unit at UCL
90 High Holborn, 2nd Floor, London WC1V 6LJ
UK
(ORCID if desired)
m.quartagno@ucl.ac.uk

Suzie M Cro
Imperial Clinical Trials Unit
Imperial College London, 1st Floor, Stadium House London, W12 7RH
UK
(ORCID if desired)
s.cro@imperial.ac.uk

James Carpenter
MRC Clinical Trials Unit at UCL
90 High Holborn, 2nd Floor, London WC1V 6LJ
UK
(ORCID if desired)
j.carpenter@ucl.ac.uk